

# DWMA: Autonomous Sensorimotor Structure Acquisition, Multi-Schema Retention, and Active Hypothesis Discrimination in an Embodied World Model

Preprint — Version 2.1 (Final Research Draft) • September 2026

Anonymous Author(s)

*Affiliation Omitted for Double-Blind Review*  
September 2026

## ABSTRACT

We introduce the **Developmental World-Model Agent (DWMA)**, an embodied computational architecture designed to acquire, maintain, and adaptively update predictive representations of interactive continuous physical environments. Rather than relying on Large Language Models or static offline training corpuses, DWMA constructs grounded world models purely through online sensorimotor contingency, refference cancellation, and epistemic active inference. The agent operates without privileged access to latent physical parameters (mass, friction, restitution, or gravity), requiring these to be inferred directly from interaction residuals. We report an audited evaluation across  $N = 50$  deterministic seeds with  $B = 10,000$  bootstrap resamples: (1) 7 baseline benchmarks reproduced (8.8× forecasting MSE reduction over persistence, Cohen's  $d = 11.45$ ); (2)  $A \rightarrow B \rightarrow C$  cross-world transfer and rapid online adaptation ( $T_\epsilon = 16.32 \pm 0.62$  steps, 50/50 pass); (3) continual multi-schema retention ( $E_{A1} = 0.00500 \rightarrow E_{A2} = 0.00501$ ,  $R = 0.9901 \pm 0.0141$ , 50/50 pass vs. overwriter catastrophic forgetting  $R = -72.97$ ); (4) active structural hypothesis discrimination (Kaplan-Meier median latency of 7.0 steps with 0% censoring vs. 17.0 steps with 22.0% right-censoring [11/50] for random babbling; log-rank  $\chi^2 = 74.98$ ,  $p < 10^{-17}$  under fair decoupled RNG); and (5) systematic hyperparameter and architectural ablations demonstrating that targeted epistemic action selection, directional EMA smoothing ( $\gamma = 0.80$ ), and discrete schema memory each contribute substantially to empirical sample efficiency, noise immunity, and multi-domain retention, preventing specific failure modes (noise-induced chatter, sluggish adaptation, sample inefficiency, and catastrophic forgetting). Complete source code, configuration files, 50-seed deterministic trajectories, and turnkey reproduction scripts (`reproduce_all.py`) are archived in the repository for full independent replication.

## 1. INTRODUCTION

Contemporary AI relies heavily on autoregressive sequence modeling over symbolic tokens. While effective for textual generation, tokens lack direct causal grounding to physical dynamics or motor refference (Harnad, 1990; Brooks, 1991; Bender & Koller, 2020). In contrast, biological developmental cognition demonstrates that infants acquire physical intuition through active sensorimotor manipulation, refference cancellation, and uncertainty-directed play (Piaget, 1952; Von Holst & Mittelstaedt, 1950; Friston, 2010; Gopnik, 2012). DWMA investigates this developmental trajectory in a simulated dynamical environment, asking whether an unprivileged agent can autonomously construct reusable predictive forward models, actively discriminate between competing physical laws, and prevent catastrophic forgetting across domains.

## 2. ARCHITECTURE & MATHEMATICAL FORMULATION

At discrete time  $t$ , ego observation is  $\mathbf{o}_t = [x_t, y_t, v_{x,t}, v_{y,t}, \theta_t]^T \in \mathbb{R}^5$ . Latent mass, friction, and gravity are strictly hidden. Continuous motor actuation  $\mathbf{a}_t \in [-1, 1]^2$  generates physical thrust  $\mathbf{F}_t = \alpha \mathbf{W}_{\text{act}} \mathbf{a}_t$  with base gain  $\alpha = 0.85$ . In the Python empirical core, parameters form decoupled diagonal matrices  $\hat{\mathbf{W}} = \text{diag}(\alpha_x, \alpha_y)$  and  $\hat{\mathbf{M}} = \text{diag}(\mu_x, \mu_y)$  enabling independent orthogonal axis adaptation, while the 60 FPS interactive web visualizer utilizes an isotropic scalar formulation  $(\alpha, \mu)$  for browser rendering efficiency; both share the identical reafference cancellation and polarity inversion dynamics. Forward kinematics prediction generates error residuals, distinguishing between instantaneous Euclidean tracking error  $E_{L2}(t) = \sqrt{\|\mathbf{e}_x(t)\|^2 + \|\mathbf{e}_v(t)\|^2}$  (in Phases 2 and 3) and multi-step forecasting mean squared error (MSE in Benchmarks 1, 3, and 6). Discrete benchmarks operate with  $\text{dt} = 1.0$ , while continuous drag integrates at  $\text{dt} = 0.02 \text{ s}$  (50 Hz). Benchmark 4 active inquiry information gain ( $5.14 \pm 0.23 \text{ nats}$ ) reflects empirical Kalman-like recursive variance contraction yielding differential Gaussian entropy reduction  $\Delta H = 0.5 * \ln(\sigma_{\text{prior}}^2 / \sigma_{\text{post}}^2)$ . Benchmark 5 evaluates thresholded kinematic affordance categorization into proto-concepts (static, movable, reactive) rather than non-parametric causal DAG discovery.




Figure 1: Closed-loop architecture integrating motor actuation, reafference cancellation, predictive forward modeling, epistemic active inference, and multi-schema memory.

## Algorithm 1: Online Recursive Adaptation & Polarity Inversion

Algorithm 1: DWMA Online State Prediction and Structural Parameter Adaptation

Input : Ego state  $s_t$ , motor action  $a_t$ , active schema parameters  $\Theta_t$   
Output : Next state prediction  $\hat{s}_{t+1}$ , updated parameters  $\Theta_{t+1}$

1. Forward Kinematics Prediction:  

$$\hat{v}_{t+1} = v_t + [\alpha_{\hat{t}} * W_{\hat{t}} * a_t - v_t * (1 - \mu_{\hat{t}})] * dt$$

$$\hat{x}_{t+1} = x_t + \hat{v}_{t+1} * dt$$
2. Motor Execution & Observation:  
Execute  $a_t$  in simulator; observe true state  $s_{t+1} = [x_{t+1}, v_{t+1}]$   
Prediction error:  $e_v = v_{t+1} - \hat{v}_{t+1}$
3. Reafference Cancellation:  

$$a_{\text{motor}} = (v_{t+1} - v_t)/dt + v_t * (1 - \mu_{\hat{t}})$$
4. Directional Cosine Alignment:  

$$\cos(\theta) = \text{dot}(a_t, a_{\text{motor}}) / \max(\epsilon, ||a_t|| * ||a_{\text{motor}}||)$$
5. Evidence Accumulation:  

$$\Lambda_t = 0.80 * \Lambda_{t-1} + 0.20 * \cos(\theta)$$
6. Structural Polarity Inversion:  
if  $\Lambda_t < -0.50$  then:  

$$W_{\hat{t+1}} = -W_{\hat{t}}$$

$$\Lambda_t = 0.0$$
else:  

$$W_{\hat{t+1}} = W_{\hat{t}}$$
7. Recursive Online Parameter Updates:  

$$\mu_{\hat{t+1}} = \mu_{\hat{t}} + \eta * e_v * v_t$$

$$\alpha_{\hat{t+1}} = \alpha_{\hat{t}} + \eta * e_v * (W_{\hat{t}} * a_t)$$

## 3. CROSS-WORLD ADAPTATION (A $\rightarrow$ B $\rightarrow$ C)

Evaluated across World A ( $\mu = 0.88$ ,  $g = 9.8$ ), World B ( $\mu = 0.71$ ,  $W_{\text{act}} = -1$ ,  $g = 4.2$ ), and World C ( $\mu = 0.12$ ,  $W_{\text{act}} = 0.60$   $\Rightarrow \alpha_{\text{eff}} = 0.51$ ,  $g = 15.7$ ). Across the initial post-transfer window, mean tracking shocks were  $1.3079 \pm 0.0421$  in World B and  $0.8117 \pm 0.0272$  in World C (Table 1). At instantaneous entry ( $t = 1$ ), unadapted inverted thrust generated an initial peak shock of  $E(0) = 1.8491 \pm 0.3356$ . Online interaction in World B restored tracking error below  $\epsilon = 0.20$  within  $T_{\epsilon} = 16.32 \pm 0.62$  steps (50/50 pass).




Figure 2: Representative single-seed empirical error trajectory in World B (Seed 1) showing instantaneous entry shock at  $t=1$  ( $E(0) = 1.8491$ ), autonomous polarity inversion at step 7.0, and error collapse ( $T_{\epsilon} = 16.32 \pm 0.62$  steps).

## 4. CONTINUAL RETENTION (A → B → A)

Multi-schema memory isolates distinct physical worlds, achieving 99.01% retention upon re-entry ( $E_{A1} = 0.00500 \rightarrow E_{A2} = 0.00501$ ), completely avoiding the catastrophic forgetting of single-model overwriting ( $E_{A2} = 0.37460$ ,  $R = -72.97$ ).




Figure 3: Retention error comparison confirming schema preservation in Full DWMA vs catastrophic forgetting in overwriter baseline.

## 5. NONLINEAR STRUCTURAL HYPOTHESIS DISCRIMINATION

When distinguishing between linear drag (H1) and quadratic drag (H2), model evidence is determined via the precision-weighted Gaussian log-likelihood ratio proxy  $\ln B21 = (SSE1 - SSE2)/(2 \cdot \sigma^2)$  ( $\sigma = 0.05$ ) under flat model priors and equivalent parameter support. Under fair decoupled RNG between environment sensor noise and agent motor selection, epistemic action selection targeting maximal discriminative divergence  $\Delta(v) = |F_{\hat{1}}(v) - F_{\hat{2}}(v)|$  achieved decisive separation in a Kaplan-Meier median of **7.0 steps** (0% censored) vs **17.0 steps** for random babbling (22.0% right-censored at  $T=30$  [11/50]; log-rank  $\chi^2 = 74.98$ ,  $p < 10^{-17}$ ). This corresponds to a **2.43× reduction in median latency** (17.0 / 7.0 steps); among uncensored successful runs, mean latency was  $7.18 \pm 2.59$  versus  $15.28 \pm 5.88$  steps (a **2.13× speedup**).




Figure 4: Kaplan-Meier survival analysis for hypothesis discrimination latency under fair decoupled RNG. Log-rank  $\chi^2 = 74.98$  ( $p < 10^{-17}$ ; 22% right-censored [11/50]; 2.43× median reduction, 2.13× uncensored speedup).

6. CONSOLIDATED MASTER EMPIRICAL EVALUATION MATRIX

Table 1 presents the complete performance profile across all 7 frozen baseline benchmarks and the 3 live empirical stress-tests. All values represent mean ± sample standard deviation across N = 50 deterministic seeds, alongside both parametric and non-parametric bootstrap 95% confidence intervals (B = 10,000 resamples).

Table 1: Consolidated Master Empirical Evaluation Matrix across N = 50 deterministic seeds (B = 10,000 bootstrap iterations).

| Phase & Phenomenon  | Condition / Policy | Primary Metric      | Observed (Mean ± SD) | Parametric 95% CI   | Bootstrap 95% CI                  | Empirical Pass Rate     |
|---------------------|--------------------|---------------------|----------------------|---------------------|-----------------------------------|-------------------------|
| Phase 1: BM-1       | Zero-Shot Transfer | Adaptation Gain     | 85.82% ± 2.53%       | [85.12, 86.53]%     | [85.15, 86.51]%                   | 49/50 (98.0%) [PASS]    |
| Phase 1: BM-2       | Blind Navigation   | Steps to Goal       | 23.10 ± 2.46 st      | [22.42, 23.78]      | [22.44, 23.78]                    | 50/50 (100.0%) [PASS]   |
| Phase 1: BM-3       | Silent Drift       | Shock MSE           | 1.384 ± 0.120        | [1.351, 1.417]      | [1.351, 1.416]                    | 50/50 (100.0%) [PASS]   |
| Phase 1: BM-4       | Active Inquiry     | Info Gain (ΔH nats) | 5.14 ± 0.23 nats     | [5.08, 5.21]        | [5.08, 5.21]                      | 50/50 (100.0%) [PASS]   |
| Phase 1: BM-5       | Causal Repair      | Affordance Ratio    | 1.18 ± 0.48 px       | [1.05, 1.32]        | [1.05, 1.32]                      | 50/50 (100.0%) [PASS]   |
| Phase 1: BM-6       | Composition        | Composition MSE     | 0.119 ± 0.075        | [0.098, 0.139]      | [0.098, 0.139]                    | 50/50 (100.0%) [PASS]   |
| Phase 1: BM-7       | Polarity Inversion | Alignment cos(θ)    | 0.9915 ± 0.0057      | [0.9899, 0.9931]    | [0.9899, 0.9931]                  | 48/50 (96.0%) [PASS]    |
| Phase 2: Baseline   | Tracking A → A     | Error E_L2(t)       | 0.0050 ± 0.0000      | [0.0050, 0.0050]    | [0.0050, 0.0050]                  | 50/50 (100.0%) [PASS]   |
| Phase 2: Transfer   | Shock A → B        | Mean Shock E_L2     | 1.3079 ± 0.0421      | [1.2963, 1.3196]    | [1.2968, 1.3199]                  | Transfer Shock          |
| Phase 2: Transfer   | Shock A → C        | Mean Shock E_L2     | 0.8117 ± 0.0272      | [0.8042, 0.8192]    | [0.8041, 0.8192]                  | Transfer Shock          |
| Phase 2: Adaptation | Online Re-tuning   | Latency T_ε         | 16.32 ± 0.62 st      | [16.15, 16.49]      | [16.14, 16.48]                    | 50/50 (100.0%) [PASS]   |
| Phase 2: Adaptation | Online Re-tuning   | Rate A_adapt        | 0.0461 ± 0.0084      | [0.0438, 0.0485]    | [0.0438, 0.0484]                  | 50/50 (100.0%) [PASS]   |
| Phase 3: Retention  | Full DWMA          | E_A1 → E_A2         | 0.00500 → 0.00501    | [0.00500, 0.00502]  | [0.00500, 0.00502]                | 50/50 (R=0.9901) [PASS] |
| Phase 3: Retention  | Overwriter         | E_A1 → E_A2         | 0.00500 → 0.37460    | [0.3657, 0.3835]    | [0.3657, 0.3835]                  | 0/50 (Catastrophic)     |
| Phase 4: Epistemic  | Full Epistemic     | Proxy ln(B21)       | 30.35 ± 8.03         | [28.12, 32.57]      | [28.10, 32.56]                    | 50/50 (100.0%) [PASS]   |
| Phase 4: Babbling   | Random Motor Play  | Proxy ln(B21)       | 24.40 ± 18.52        | [16.03, 24.00]      | [16.01, 23.95]                    | 39/50 (78.0%) [PASS]    |
| Phase 4: Latency    | KM Median Contrast | Discrim. Steps      | 7.0 vs 17.0 st       | Mantel-Cox Log-Rank | χ² = 74.98, p < 10 <sup>-17</sup> | Decisive Speedup        |

7. ABLATION STUDIES & HYPERPARAMETER SENSITIVITY ANALYSIS

To substantiate the causal mechanisms underlying DWMA and demonstrate that performance is not an artifact of handcrafted constants, we executed an ablation suite across N = 50 seeds.

7.1 Component & Policy Ablations

Table 2: Exploration policy and architectural memory ablation matrix (N = 50 seeds).

| Architecture / Policy Condition                | Median Latency | Empirical Success Rate | Right-Censored Runs | Continual Retention R          |
|------------------------------------------------|----------------|------------------------|---------------------|--------------------------------|
| Full DWMA (Targeted Epistemic $v^*$ )          | 7.0 steps      | 50/50 (100.0%)         | 0/50 (0.0%)         | $0.9901 \pm 0.0141$            |
| Epsilon-Greedy Epistemic ( $\epsilon = 0.20$ ) | 7.0 steps      | 50/50 (100.0%)         | 0/50 (0.0%)         | $0.9901 \pm 0.0141$            |
| Random Babbling Baseline (Uniform Play)        | 17.0 steps     | 39/50 (78.0%)          | 11/50 (22.0%)       | —                              |
| Passive Observer Baseline ( $a = 0$ )          | >30.0 steps    | 0/50 (0.0%)            | 50/50 (100.0%)      | —                              |
| No Multi-Schema Memory (Overwriting)           | 7.0 steps      | 50/50 (100.0%)         | 0/50 (0.0%)         | -72.97 (Catastrophic Collapse) |

**Causal Interpretation:** Targeted epistemic exploration drives the agent directly to the velocity regime where linear drag ( $-k_1 \cdot v$ ) and quadratic drag ( $-k_2 \cdot v|v|$ ) exhibit maximal divergence. Passive observation fails completely because environmental damping without actuation leaves velocity near zero where both hypotheses make near-identical predictions.

7.2 Hyperparameter Sensitivity: Smoothing Factor  $\gamma$  and Inversion Threshold  $\Lambda$

Table 3: Sensitivity analysis of directional EMA factor  $\gamma$  and structural inversion threshold  $\Lambda_{\text{thresh}}$  across N = 50 seeds.

| Parameter Varied                                                           | Setting Tested              | Mean Inversion Step    | Trigger Rate | Mechanistic Assessment & Trade-off                |
|----------------------------------------------------------------------------|-----------------------------|------------------------|--------------|---------------------------------------------------|
| Directional EMA Smoothing $\gamma$ (at $\Lambda_{\text{thresh}} = -0.50$ ) | $\gamma = 0.60$             | $2.00 \pm 0.00$ steps  | 50/50        | Vulnerable to transient sensor noise chatter      |
|                                                                            | $\gamma = 0.70$             | $2.00 \pm 0.00$ steps  | 50/50        | Marginal noise immunity under collision spikes    |
|                                                                            | $\gamma = 0.80$ (Default)   | $4.00 \pm 0.00$ steps  | 50/50        | Optimal balance: noise-immune & fast recovery     |
|                                                                            | $\gamma = 0.90$             | $7.00 \pm 0.00$ steps  | 50/50        | Sluggish adaptation delay (prolongs high error)   |
|                                                                            | $\gamma = 0.95$             | $14.00 \pm 0.00$ steps | 50/50        | Excessive smoothing (3.5× slower polarity flip)   |
| Inversion Threshold $\Lambda_{\text{thresh}}$ (at $\gamma = 0.80$ )        | $\Lambda = -0.30$           | $2.00 \pm 0.00$ steps  | 50/50        | Premature trigger risk during sharp braking       |
|                                                                            | $\Lambda = -0.40$           | $3.00 \pm 0.00$ steps  | 50/50        | False-inversion risk during rapid decelerations   |
|                                                                            | $\Lambda = -0.50$ (Default) | $4.00 \pm 0.00$ steps  | 50/50        | Optimal: zero false positives on familiar physics |
|                                                                            | $\Lambda = -0.60$           | $5.00 \pm 0.00$ steps  | 50/50        | Delayed adaptation response                       |
|                                                                            | $\Lambda = -0.70$           | $6.00 \pm 0.00$ steps  | 50/50        | Over-conservative: delays recovery by 50%         |

## 8. CONCLUSION & OPEN REPRODUCIBILITY

DWMA v2.1 demonstrates that an unprivileged embodied agent can autonomously acquire predictive forward models, adapt to severe dynamical domain shifts, prevent catastrophic forgetting via multi-schema memory, and discriminate nonlinear physical hypotheses through epistemic active inference. Ablation experiments indicate that targeted epistemic action selection, EMA directional filtering, and discrete schema memory make measurable contributions to the evaluated outcomes, stabilizing online dynamics, accelerating hypothesis discrimination by 2.43× in median latency (2.13× on uncensored runs), and insulating prior physical models from catastrophic interference. Complete Python/JS source code, configuration files, 50-seed deterministic trajectories, and turnkey reproduction scripts (`python reproduce_all.py`) are archived in the repository for complete independent replication.

## DECLARATION OF GENERATIVE AI AND AI-ASSISTED TECHNOLOGIES

During the preparation of this work, the author utilized generative AI and AI-assisted tools (including deep research assistants and large language models) for drafting, restructuring, editorial refinement, LaTeX formatting, and assistance with Python reproduction scripts under continuous human direction and oversight. Following the use of these tools, the author independently audited, executed, validated, and edited all empirical code, mathematical formulations, figure visualizations, and manuscript text. The author assumes full personal responsibility for the factual validity, scientific integrity, and conclusions presented in this publication.

## REFERENCES

35 foundational and contemporary references spanning world models, active inference, predictive coding, symbol grounding, and continual learning.

1. Barsalou, L. W. (1999). Perceptual symbol systems. *Behavioral and Brain Sciences*, 22(4), 577-660.
2. Bellemare, M., et al. (2016). Unifying count-based exploration and intrinsic motivation. *NeurIPS*, 29.
3. Bender, E. M., & Koller, A. (2020). Climbing towards NLU: On meaning, form, and understanding. *ACL*, 5185-5198.
4. Bisk, Y., et al. (2020). Experience grounds language. *EMNLP*, 8718-8735.
5. Bogacz, R. (2017). A tutorial on the free-energy framework for modelling perception and learning. *J. Math. Psychol.*, 76, 198-211.
6. Brooks, R. A. (1991). Intelligence without representation. *Artificial Intelligence*, 47(1-3), 139-159.
7. Buckley, C. L., et al. (2021). The free energy principle for action and perception: A mathematical review. *J. Math. Psychol.*, 86, 102558.
8. Burda, Y., et al. (2018). Exploration by random network distillation. *ICLR*.
9. Cangelosi, A. (2010). Grounding language in action and perception: From cognitive agents to humanoid robots. *Phys. Life Rev.*, 7(2), 139-151.
10. Clark, A. (2013). Whatever next? Predictive brains, situated agents, and the future of cognitive science. *Behav. Brain Sci.*, 36(3), 181-204.
11. French, R. M. (1999). Catastrophic forgetting in connectionist networks. *Trends in Cognitive Sciences*, 3(4), 128-135.
12. Friston, K. (2005). A theory of cortical responses. *Phil. Trans. R. Soc. B*, 360(1456), 815-836.
13. Friston, K. (2010). The free-energy principle: A unified brain theory? *Nature Reviews Neuroscience*, 11(2), 127-138.
14. Friston, K., et al. (2015). Active inference and epistemic value. *Cognitive Neuroscience*, 6(4), 187-214.
15. Friston, K., et al. (2017). Active inference, curiosity and insight. *Neural Computation*, 29(10), 2633-2683.
16. Gopnik, A. (2012). Scientific thinking in young children: Theoretical advances, empirical research, and policy implications. *Science*, 337(6102), 1623-1627.
17. Ha, D., & Schmidhuber, J. (2018). Recurrent world models facilitate policy evolution. *NeurIPS*, 31.
18. Hafner, D., et al. (2020). Dream to control: Learning behaviors by latent imagination. *ICLR*.
19. Hafner, D., et al. (2023). Mastering diverse domains through world models. *arXiv:2301.04104*.
20. Harnad, S. (1990). The symbol grounding problem. *Physica D: Nonlinear Phenomena*, 42(1-3), 335-346.
21. Kaplan, E. L., & Meier, P. (1958). Nonparametric estimation from incomplete observations. *JASA*, 53(282), 457-481.
22. Kirkpatrick, J., et al. (2017). Overcoming catastrophic forgetting in neural networks. *PNAS*, 114(13), 3521-3526.
23. LeCun, Y. (2022). A path towards autonomous machine intelligence. *Open Review*.
24. Mantel, N. (1966). Evaluation of survival data and two new rank order statistics. *Cancer Chemother. Rep.*, 50(3), 163-170.
25. McCloskey, M., & Cohen, N. J. (1989). Catastrophic interference in connectionist networks. *Psychol. Learn. Motiv.*, 24, 109-165.
26. Millidge, B., et al. (2021). Whence the expected free energy? *Neural Computation*, 33(2), 447-482.
27. Oudeyer, P. Y., & Kaplan, F. (2007). What is intrinsic motivation? A typology of computational approaches. *Front. Neurobot.*, 1, 6.
28. Pathak, D., et al. (2017). Curiosity-driven exploration by self-supervised prediction. *ICML*, 2778-2787.
29. Piaget, J. (1952). *The origins of intelligence in children*. International Universities Press.
30. Rao, R. P., & Ballard, D. H. (1999). Predictive coding in the visual cortex. *Nature Neuroscience*, 2(1), 79-87.
31. Rolnick, D., et al. (2019). Experience replay for continual learning. *NeurIPS*, 32.
32. Rusu, A. A., et al. (2016). Progressive neural networks. *arXiv:1606.04671*.
33. Schmidhuber, J. (1990). Making the world differentiable: On combining models of the environment with reinforcement learning. *TR FKI-126-90*.
34. Von Holst, E., & Mittelstaedt, H. (1950). Das Reafferenzprinzip. *Naturwissenschaften*, 37(20), 464-476.
35. Whittington, J. C., et al. (2020). The Tolman-Eichenbaum machine: Unifying space and relational memory in the hippocampal formation. *Cell*, 183(5), 1249-1263.